


EXP1 Figure 1 (BIG_GPU_FINAL): InfluCoder vs gradient-influence baselines -- cost/quality, scaling, and amortization


A

EXP1 Part 1 (BIG_GPU_FINAL) — 400x400 eval
LESS r=32, LoGRA r=32 (uniform, batched)



B

EXP1 Part 2 (BIG_GPU_FINAL) — InfluCoder scaling, 400x400 eval
epochs=8, hard_ratio=0.0, lr=5e-05



C

EXP1 Part 3 (BIG_GPU_FINAL) — GPU-time vs samples-processed amortization
InfluCoder rate measured directly through n=100,000 (not extrapolated)

